

9v6okdfz3

March 10, 2025

```
[1]: #!pip install kagglehub#!pip install matplotlib  
#!pip install -U scikit-learn  
#!pip install matplotlib  
#!pip install -U scikit-learn  
#!pip install numpy  
#!pip install pandas  
#!pip install scikit-image  
#!pip install imblearn
```

```
[2]: import os  
import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
from tqdm import tqdm  
from skimage import io, transform, feature  
from sklearn.model_selection import train_test_split, StratifiedKFold  
from sklearn.svm import SVC  
from sklearn.metrics import classification_report, confusion_matrix  
from sklearn.utils.multiclass import unique_labels  
from imblearn.over_sampling import SMOTE  
from sklearn.preprocessing import LabelEncoder  
from sklearn.model_selection import GridSearchCV  
  
# Configurations  
dataset_root = "/Users/markus/.cache/kagglehub/datasets/pulavendrancelvaraj/  
↳oasis-dataset/versions/1/input"  
classes = ["Non Demented", "Very mild Dementia", "Mild Dementia", "Moderate_  
↳Dementia"]
```

```
[3]: # Chargement des données  
X, y = [], []  
for class_name in classes:  
    class_path = os.path.join(dataset_root, class_name)  
    for filename in tqdm(os.listdir(class_path), desc=f"Loading {class_name}"):   
        try:  
            img_path = os.path.join(class_path, filename)  
            img = io.imread(img_path, as_gray=True)
```

```

        img = transform.resize(img, (64, 64), anti_aliasing=True)
    ↪ #Redimensionnement
        img = img / 255.0 #Normalisation
        fd = feature.hog(img, orientations=8, pixels_per_cell=(16, 16),
                        cells_per_block=(1, 1), channel_axis=None)
    ↪ #Extraction des caractéristiques HOG
        X.append(fd)
        y.append(class_name)
    except Exception as e:
        print(f"Error loading {img_path}: {e}")

```

```

Loading Non Demented: 100%|          | 3000/3000 [00:38<00:00, 78.68it/s]
Loading Very mild Dementia: 100%|    | 3000/3000 [00:37<00:00, 79.10it/s]
Loading Mild Dementia: 100%|         | 3000/3000 [00:38<00:00, 77.54it/s]
Loading Moderate Dementia: 100%|     | 488/488 [00:06<00:00, 77.21it/s]

```

```

[4]: # Vérification de classes chargées
print(f"\n Données chargées")
print(f"Labels uniques : {np.unique(y)}")
print(f"Nombre d'échantillons : {len(X)}")

```

```

Données chargées
Labels uniques : ['Mild Dementia' 'Moderate Dementia' 'Non Demented' 'Very mild
Dementia']
Nombre d'échantillons : 9488

```

```

[5]: # Chemin de l'image testée
img_path = "/Users/markus/.cache/kagglehub/datasets/pulavendrancelvaraj/
    ↪ oasis-dataset/versions/1/input/Mild Dementia/OAS1_0028_MR1_mpr-1_103.jpg"

# Chargement de l'image en niveaux de gris
img = io.imread(img_path, as_gray=True)

# Redimensionnement à 64x64 avec anti-aliasing
img_resized = transform.resize(img, (64, 64), anti_aliasing=True)

# Normalisation
img_normalized = img_resized / 255.0

# Extraction des features HOG et visualisation
fd, hog_image = feature.hog(img_normalized, orientations=8,
    ↪ pixels_per_cell=(16, 16),
                                cells_per_block=(1, 1), visualize=True,
    ↪ channel_axis=None)

```

```
[6]: # Affichage des résultats
fig, ((ax1, ax2), (ax3, ax4)) = plt.subplots(2, 2, figsize=(12, 8))


# Image originale
ax1.imshow(img, cmap='gray')
ax1.set_title("Image originale", fontsize=14)
ax1.axis('off')

# Image redimensionnée
ax2.imshow(img_resized, cmap='gray')
ax2.set_title("Image redimensionnée (64x64)", fontsize=14)
ax2.axis('off')

# Image normalisée
ax3.imshow(img_normalized, cmap='gray')
ax3.set_title("Image normalisée", fontsize=14)
ax3.axis('off')

# Visualisation HOG
ax4.imshow(hog_image, cmap='gray')
ax4.set_title("Features HOG", fontsize=14)
ax4.axis('off')


plt.tight_layout()
plt.show()
```



```
[7]: # Analyse de la distribution des classes
unique_classes, class_counts = np.unique(y, return_counts=True)
print("\n Distribution des classes ")
for cls, count in zip(unique_classes, class_counts):
    print(f"Classe: {cls} | Nombre d'exemples: {count}")

# Visualisation
plt.figure(figsize=(8,5))
plt.bar(unique_classes, class_counts, color=['green','blue','red','yellow'])
plt.xlabel('Classe')
plt.ylabel('Nombre d\'images')
plt.title('Distribution des classes dans le dataset OASIS')
plt.show()
```

```
Distribution des classes
Classe: Mild Dementia | Nombre d'exemples: 3000
Classe: Moderate Dementia | Nombre d'exemples: 488
Classe: Non Demented | Nombre d'exemples: 3000
Classe: Very mild Dementia | Nombre d'exemples: 3000
```



```
[8]: from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import cross_val_score
```

```
[9]: # Prétraitement pour le modèle
X = np.array(X)
y = np.array(y)
label_encoder = LabelEncoder()
y_encoded = label_encoder.fit_transform(y)

# Train-Test split (sans SMOTE)
X_train, X_test, y_train, y_test = train_test_split(X, y_encoded, test_size=0.
↪2, stratify=y_encoded)
```

```
[10]: # Comparaison des modèles (sans SMOTE)
print("\n Comparaison des modèles")
models = [
    ('LR', LogisticRegression(max_iter=2000)),
    ('KNN', KNeighborsClassifier()),
    ('DT', DecisionTreeClassifier()),
    ('RF', RandomForestClassifier(n_estimators=100)),
```

```

    ('SVM', SVC(kernel='linear', C=1.0, probability=True))
]

results = []
names = []
kfold = StratifiedKFold(n_splits=10, shuffle=True)

for name, model in models:
    cv_results = cross_val_score(model, X_train, y_train, cv=kfold,
    ↪scoring='f1_macro')
    results.append(cv_results)
    names.append(name)
    print(f"{name}: {cv_results.mean():.3f} ({cv_results.std():.3f})")

```

Comparaison des modèles

LR: 0.660 (0.016)

KNN: 0.970 (0.010)

DT: 0.729 (0.022)

RF: 0.964 (0.010)

SVM: 0.702 (0.024)

```

[11]: # Baseline SVM
svm_base = SVC(kernel='linear', C=1.0, probability=True)
svm_base.fit(X_train, y_train)
y_pred_base = svm_base.predict(X_test)
print("\n Rapport de Classification (Sans Rééquilibrage)")
print(classification_report(y_test, y_pred_base, target_names=label_encoder.
    ↪classes_))

```

Rapport de Classification (Sans Rééquilibrage)

	precision	recall	f1-score	support
Mild Dementia	0.67	0.77	0.71	600
Moderate Dementia	0.81	0.62	0.71	98
Non Demented	0.73	0.73	0.73	600
Very mild Dementia	0.79	0.71	0.75	600
accuracy			0.73	1898
macro avg	0.75	0.71	0.72	1898
weighted avg	0.74	0.73	0.73	1898

```


[12]: # Matrice de confusion
cm = confusion_matrix(y_test, y_pred_base)
plt.figure(figsize=(8,6))

```

```

plt.imshow(cm, interpolation='nearest', cmap='Blues')
plt.title("Matrice de Confusion (Sans SMOTE)")
plt.colorbar()
tick_marks = np.arange(len(classes))
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)
plt.tight_layout()
plt.ylabel('Classe réelle')
plt.xlabel('Classe prédite')
plt.show()

```



```

[13]: # Rééquilibrage avec SMOTE
smote = SMOTE()
X_train_res, y_train_res = smote.fit_resample(X_train, y_train)

# Rééquilibrage
_, counts = np.unique(y_train_res, return_counts=True)

```

```
print("\n Classes après SMOTE ")
for cls, cnt in zip(unique_classes, counts):
    print(f"{cls}: {cnt}")
```

```
Classes après SMOTE
Mild Dementia: 2400
Moderate Dementia: 2400
Non Demented: 2400
Very mild Dementia: 2400
```

```
[14]: # GridSearchCV
param_grid = {
    'C': [0.1, 1, 10],
    'kernel': ['linear', 'rbf'],
    'gamma': ['scale', 0.01, 0.1]}
```

```
[15]: # Cross-validation
cv = StratifiedKFold(n_splits=3, shuffle=True)
grid_search = GridSearchCV(SVC(), param_grid, cv=cv, n_jobs=-1,
    ↳scoring='f1_macro')
```

```
[16]: # Entraînement avec GridSearch sur données SMOTED
grid_search.fit(X_train_res, y_train_res)

print(f"Meilleurs paramètres: {grid_search.best_params}")
```

```
Meilleurs paramètres: {'C': 10, 'gamma': 'scale', 'kernel': 'rbf'}
```

```
[17]: # Evaluation du modèle optimisé
best_svm = grid_search.best_estimator_
y_pred = best_svm.predict(X_test)
```

```
[18]: # Rapport de classification final
print("\n Rapport de Classification (avec SMOTE + Optimisation) ")
print(classification_report(y_test, y_pred, target_names=label_encoder.
    ↳classes_))
```

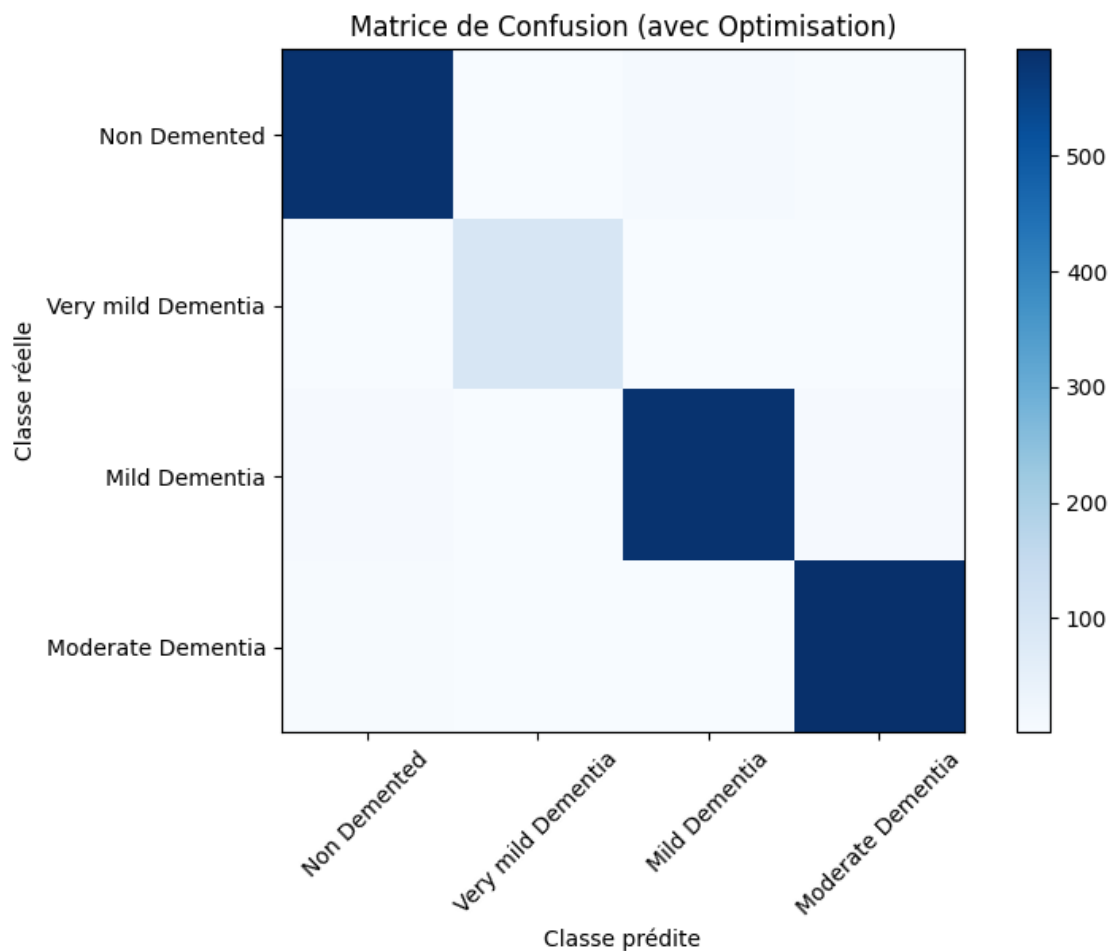
```
Rapport de Classification (avec SMOTE + Optimisation)
              precision    recall  f1-score   support

Mild Dementia      0.98      0.98      0.98        600
Moderate Dementia  0.97      0.96      0.96         98
Non Demented       0.98      0.97      0.98        600
Very mild Dementia 0.98      0.99      0.98        600

accuracy                   0.98      1898
```


macro avg	0.98	0.97	0.98	1898
weighted avg	0.98	0.98	0.98	1898

```
[19]: # Matrice de confusion finale
cm_optimized = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(8,6))
plt.imshow(cm_optimized, interpolation='nearest', cmap='Blues')
plt.title("Matrice de Confusion (avec Optimisation)")
plt.colorbar()
plt.xticks(tick_marks, classes, rotation=45)
plt.yticks(tick_marks, classes)
plt.tight_layout()
plt.ylabel('Classe réelle')
plt.xlabel('Classe prédite')
plt.show()
```



```
[20]: # Analyse de l'impact de SMOTE
print("\n Comparaison des performances ")
print("Précision avant SMOTE:", svm_base.score(X_test, y_test))
print("Précision après SMOTE:", best_svm.score(X_test, y_test))
```

```
Comparaison des performances
Précision avant SMOTE: 0.7297154899894626
Précision après SMOTE: 0.9783983140147524
```

```
[21]: # Validation croisée sur le modèle optimisé
cv_scores = cross_val_score(best_svm, X, y_encoded, cv=cv)
print(f"\nValidation croisée (Moyenne): {cv_scores.mean():.3f} ± {cv_scores.
    ↳std():.3f}")
```

```
Validation croisée (Moyenne): 0.976 ± 0.003
```